

# Medical Insurance Cost Prediction

Presented By: Muhammad Osama



# Introduction

- ▶ **Objective:**

Predict insurance costs and classify high-cost patients.

- ▶ **Dataset:**

- ▶ Source: Kaggle
- ▶ Size: 1338 records, 7 features.

- ▶ **Features**

- ▶ Numerical: Age, BMI, Children, Charges
- ▶ Categorical: Sex, Smoker, Region

- ▶ **Importance:**

Insurance pricing struggles to reflect individual risk accurately.

Need for robust models to predict medical costs and identify high-risk individuals..

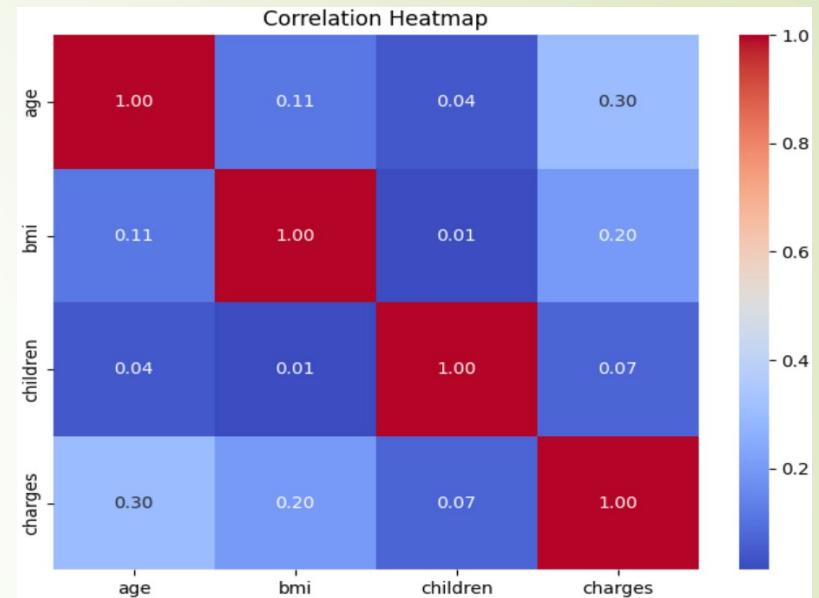
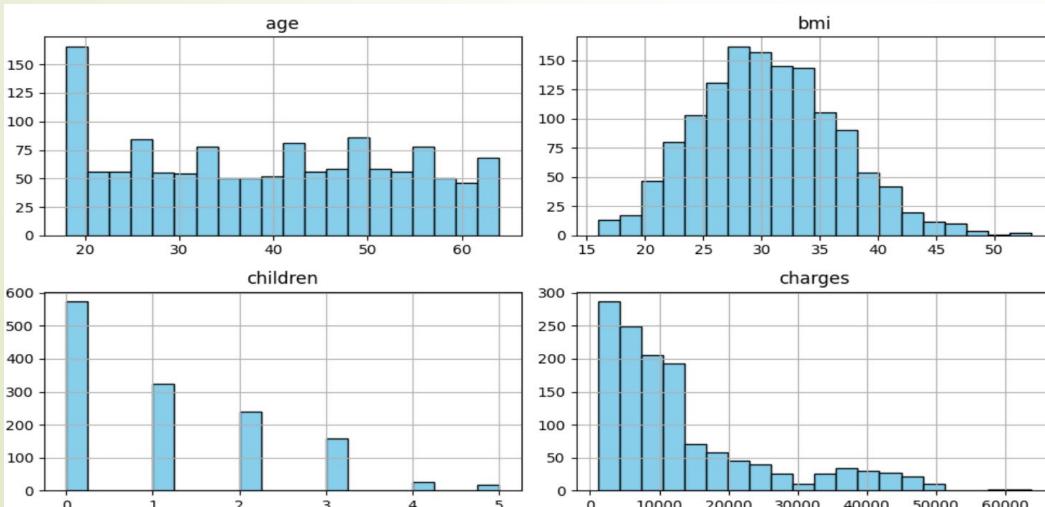
# Exploratory Data Analysis (EDA)

## ► Correlation Insights

- Age & Charges: Strongest correlation (0.30)
- BMI: Moderate impact (0.20)

## ► Distributions & Outliers:

- Age shows bimodal peaks (20s & 50s).
- Charges are right-skewed (max = \$63,770)
- Smokers pay ~ 3.8x more on average



# Data Preprocessing

- ▶ **Importing Libraries:** Importing all necessary libraries
- ▶ **Missing Values:** Check for missing values
- ▶ **Label Encoding:** Label Encoded all categorical features
- ▶ **Feature Scaling :** StandardScaler normalized input features
- ▶ **New Target:** high\_cost = charges > median (binary classification)
- ▶ **Train Test Split:** 80% training, 20% testing.

```
# Check missing values
print("Missing values:\n", df.isnull().sum())

# Encode categorical features
df_encoded = df.copy()
label_encoders = {}

for col in ['sex', 'smoker', 'region']:
    le = LabelEncoder()
    df_encoded[col] = le.fit_transform(df_encoded[col])
    label_encoders[col] = le

# Create classification target: high_cost (1 if charges > median)
median_charge = df_encoded['charges'].median()
df_encoded['high_cost'] = (df_encoded['charges'] > median_charge).astype(int)

# Features for both models
features = ['age', 'sex', 'bmi', 'children', 'smoker', 'region']

# Scale features
scaler = StandardScaler()
df_encoded[features] = scaler.fit_transform(df_encoded[features])

# Define regression and classification targets
X = df_encoded[features]
y_reg = df_encoded['charges']
y_clf = df_encoded['high_cost']

# Split data
X_train_reg, X_test_reg, y_train_reg, y_test_reg = train_test_split(X, y_reg, test_size=0.2, random_state=42)
X_train_clf, X_test_clf, y_train_clf, y_test_clf = train_test_split(X, y_clf, test_size=0.2, random_state=42)
```

# Baseline Models

## ► Linear Regression: For Cost Prediction

- **R<sup>2</sup> Score:** 0.783
- **MSE:** 33635210
- **MMAE:** 4186

## ► Logistic Regression: For High-Cost classification

- **Accuracy:** 91%
- **Precision:** 88%
- **F1 Score:** 90%



### Confusion Matrix Summary

Metric	Value	Description
True Positives (TP)	131	Correctly predicted high-cost patients
False Positives (FP)	15	Predicted high-cost, but actually low-cost
False Negatives (FN)	9	Predicted low-cost, but actually high-cost
True Negatives (TN)	113	Correctly predicted low-cost patients

# Deep Learning Models

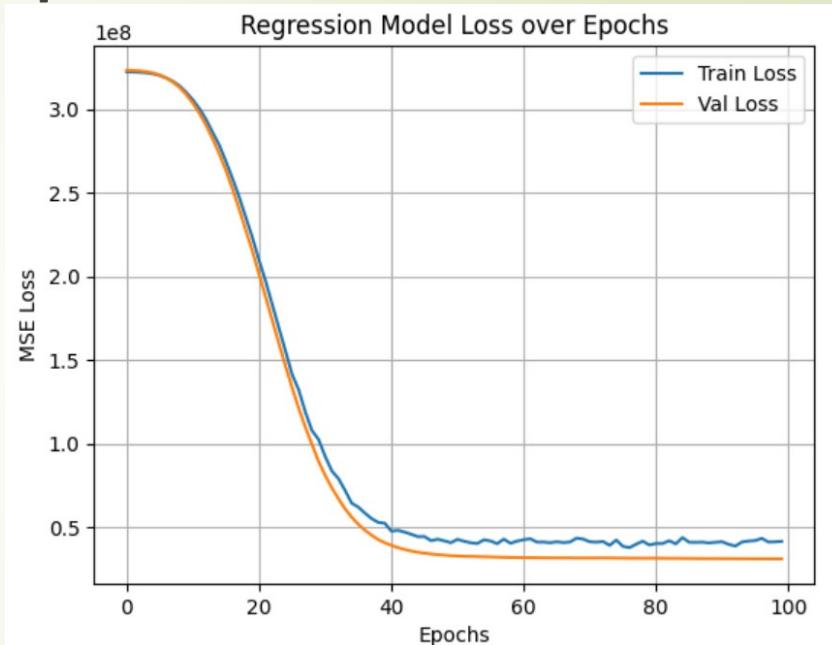
## ► Regression Model (DNN)

- **Architecture:** [64 → Dropout → 32 → Dropout → Output]
- **Activation Function:** Relu
- **Optimizer:** Adam
- **Loss:** MSE
- **Early Stopping**
- **Performance:**
  - R<sup>2</sup> Score: 0.8002
  - MSE: 31022003.67
  - MAE: 3,843

```
reg_model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train_reg.shape[1],)),
    Dropout(0.3),
    Dense(32, activation='relu'),
    Dropout(0.2),
    Dense(1) # Output layer with 1 node for regression
])

reg_model.compile(optimizer='adam', loss='mse', metrics=['mae'])

# Callback to stop early if val_loss doesn't improve
early_stop = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
```



# Deep Learning Models

## ► Classification Model (DNN)

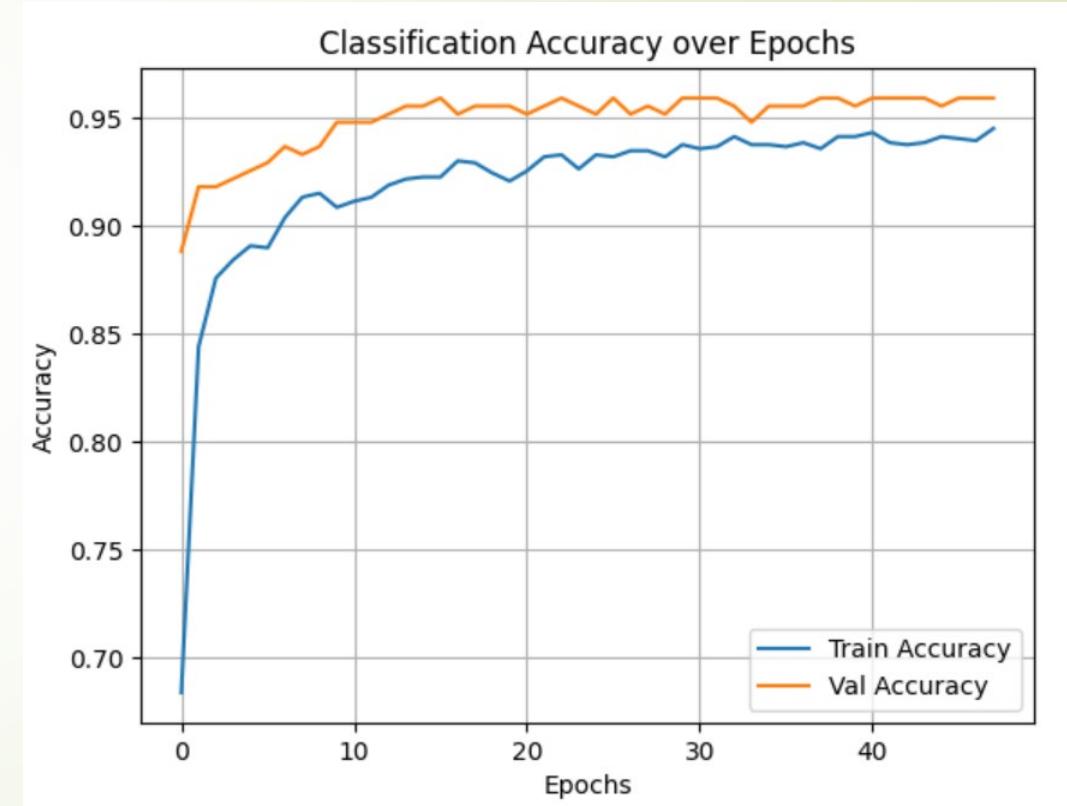
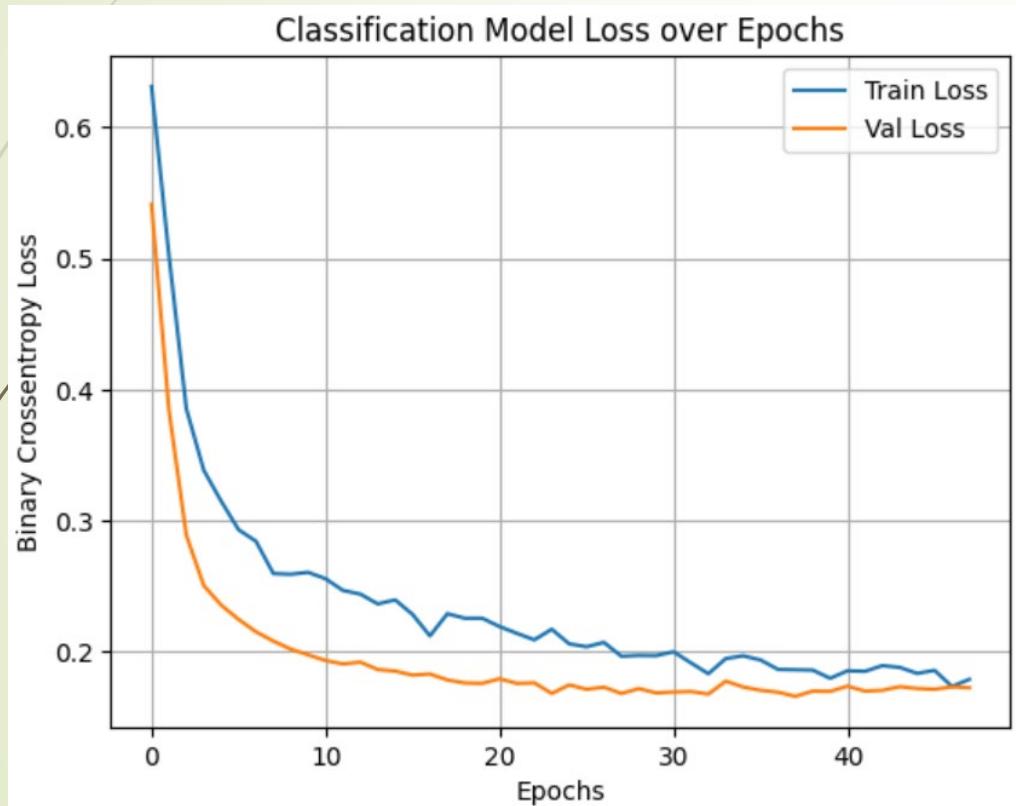
- **Architecture:** [64 → Dropout → 32 → Dropout → Sigmoid]
- **Activation Function:** Relu & Sigmoid
- **Optimizer:** Adam
- **Loss:** Binary Cross Entropy
- **Early Stopping**
- **Performance:**
  - Accuracy: 95.9 %
  - Precision: 98.3 %
  - F1 Score: 95.4 %

```
# Build model
clf_model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train_clf.shape[1],)),
    Dropout(0.3),
    Dense(32, activation='relu'),
    Dropout(0.2),
    Dense(1, activation='sigmoid') # Sigmoid for binary classification
])

clf_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

# Deep Learning Models

## ► Classification Model (DNN) Loss & Accuracy Over Epoch



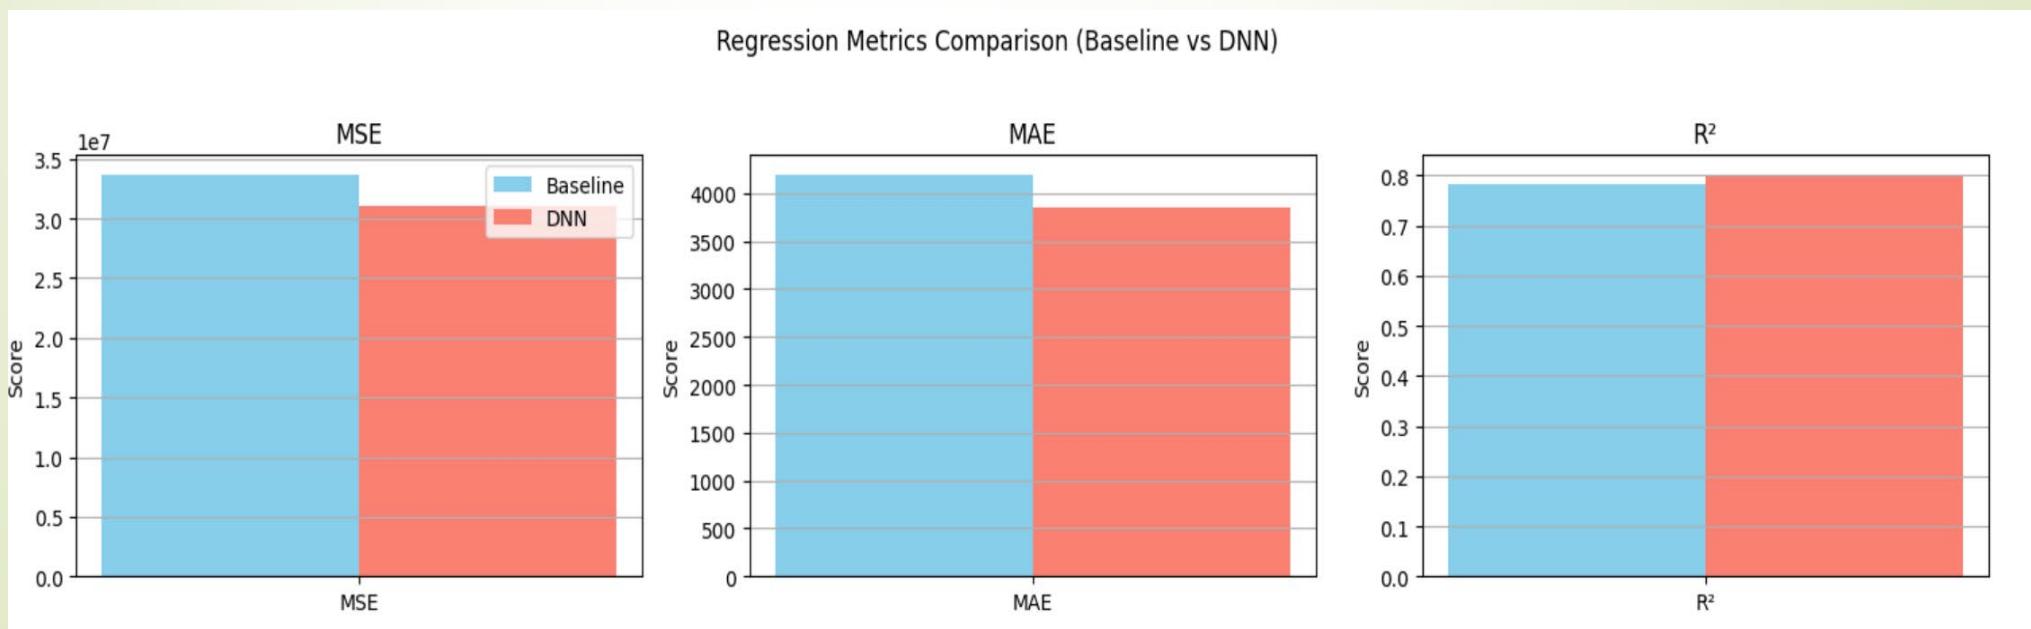
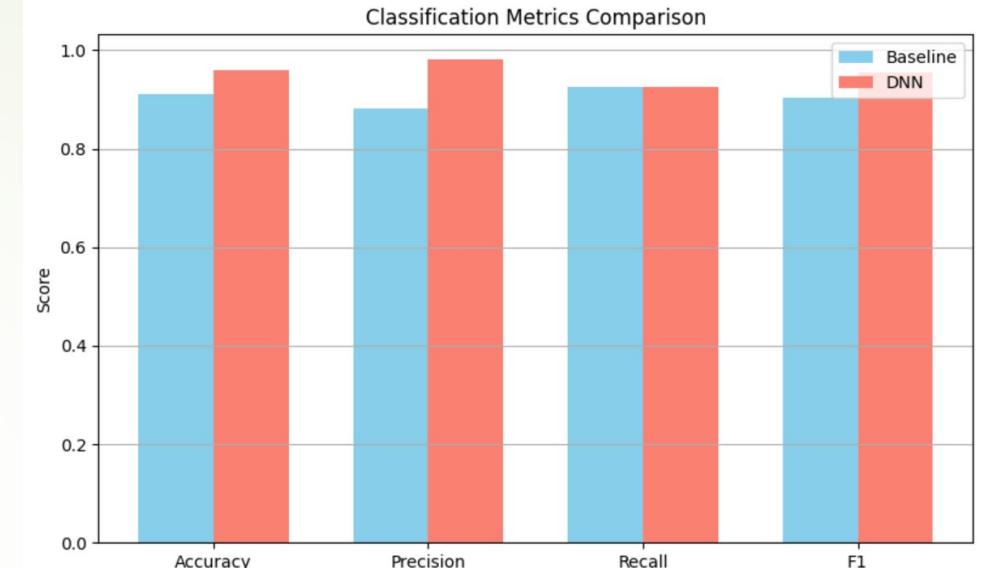
# K-Fold Cross-Validation

- ▶ A powerful method to test model stability and generalization
  - ▶ The dataset is split into K equal parts (folds)
  - ▶ The model trains on  $(K-1)$  folds and tests on the remaining one
  - ▶ This cycle repeats K times, each fold playing the test set once
  - ▶ Final performance is the average across all runs—reduces single-split
- ▶ **Results:**
- ▶ Classification DNN (5-Fold)
    - ▶ Mean Accuracy:  $94.2\% \pm 2\%$
    - ▶ Mean F1:  $94.0\% \pm 1.9\%$
  - ▶ Regression DNN (5-Fold)
    - ▶ Mean MSE:  $36.7M \pm 3.6M$
    - ▶ Mean MAE:  $4,110 \pm 144$

# Comparative Performance

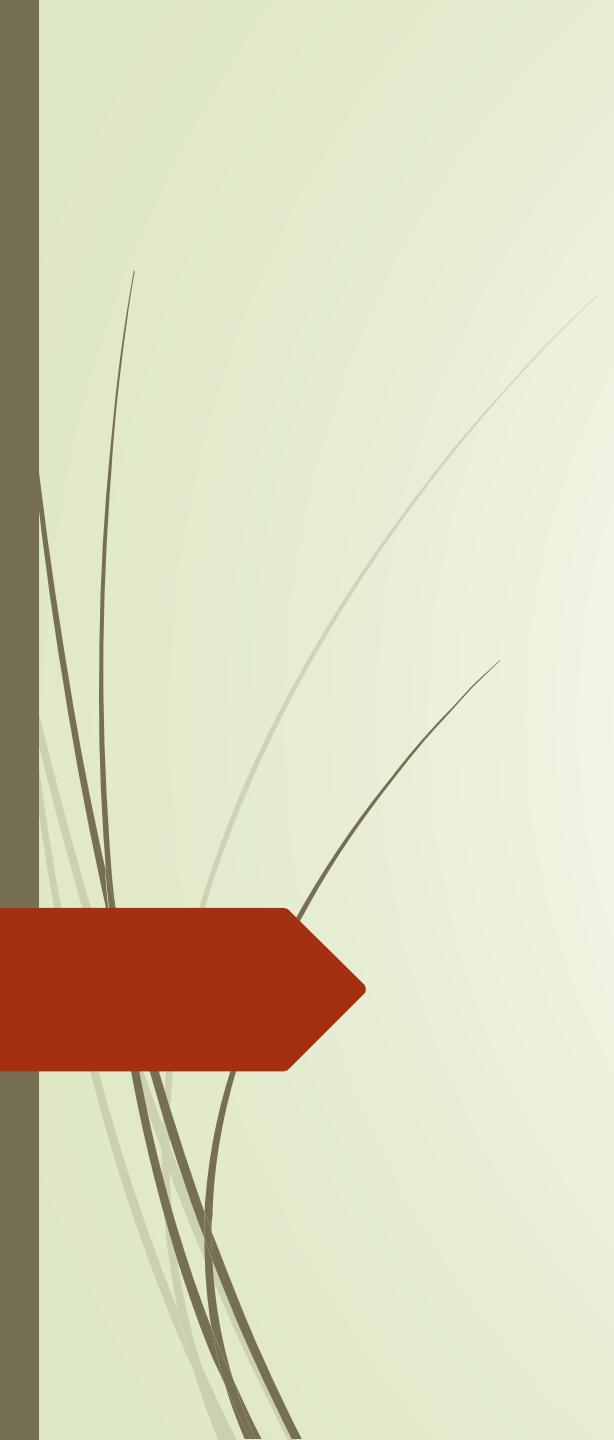
Metric	Baseline (Regression)	DNN (Regression)
MSE	33.6M	31.0M
MAE	\$4,186	\$3,843
R <sup>2</sup> Score	0.783	0.800

Metric	Baseline (Classification)	DNN (Classification)
Accuracy	91%	95.9%
Precision	88.3%	98.3%
F1 Score	90.4%	95.4%



# Conclusion

- ▶ Deep Neural Networks outperform traditional models.
- ▶ DNNs capture non-linear patterns like smoker × BMI interactions.
- ▶ Recommendations:
  - ▶ Use Bayesian optimization for hyperparameters
  - ▶ Integrate SHAP/LIME for model interpretability
  - ▶ Explore deeper or residual networks for further improvements



Thank you